


Enhancing Ethereum PoA Clique Network with DAG-based BFT Consensus

Yongrae Jo and Chanik Park

Pohang University of Science And Technology


Proof of Authority (PoA)



- Consensus for permissioned blockchain
 - Only a few messages for block agreement
 - Efficient than proof-of-work (PoW)



Proof of Authority (PoA)



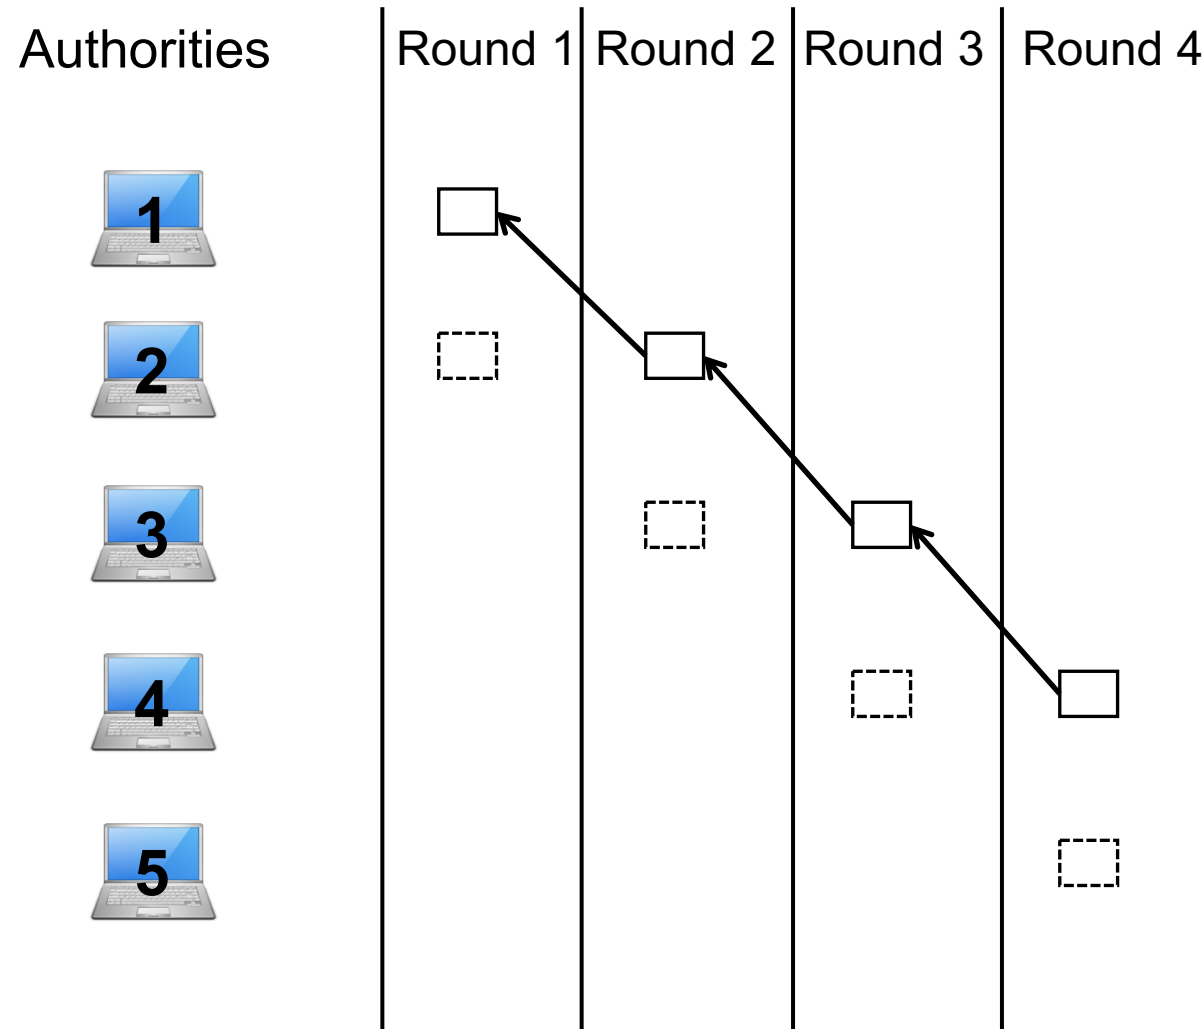
- Consensus for permissioned blockchain
 - Only a few messages for block agreement
 - Efficient than proof-of-work (PoW)



→ **Clique**
(our focus)



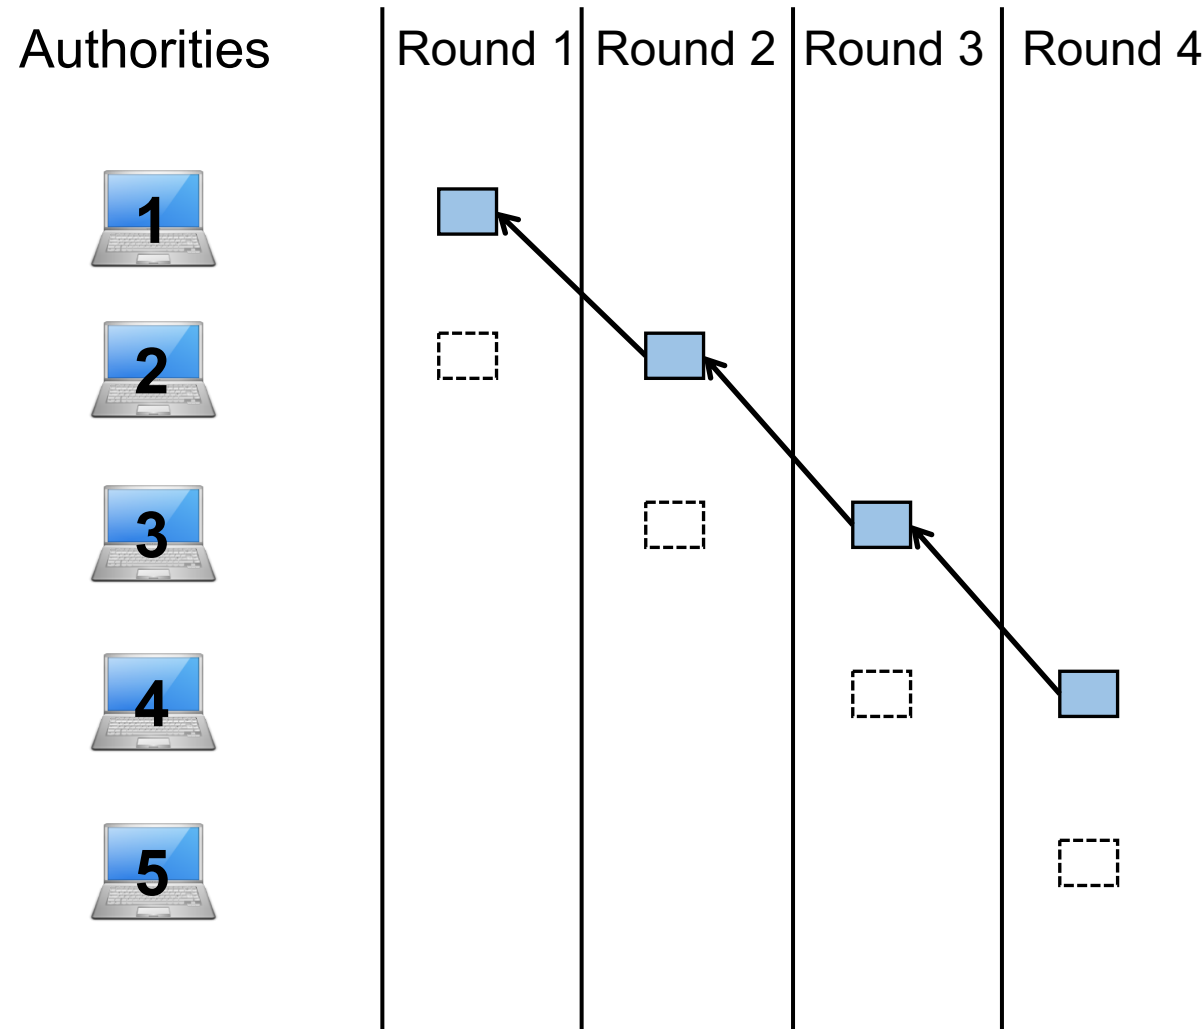
PoA Clique: An Illustration



- Round-based
 - (rotating) leader & round timeout Δ
- Multiple block proposals in a round
 - by a leader
 - by non-leaders* with random delays

*#non-leaders = up to $\left\lceil \frac{N}{2} \right\rceil - 2$


PoA Clique: An Illustration



- Round-based
 - (rotating) leader & round timeout Δ
- Multiple block proposals in a round
 - by a leader
 - by non-leaders* with random delays
- Single block acceptance in a round
 - Normal: a leader block

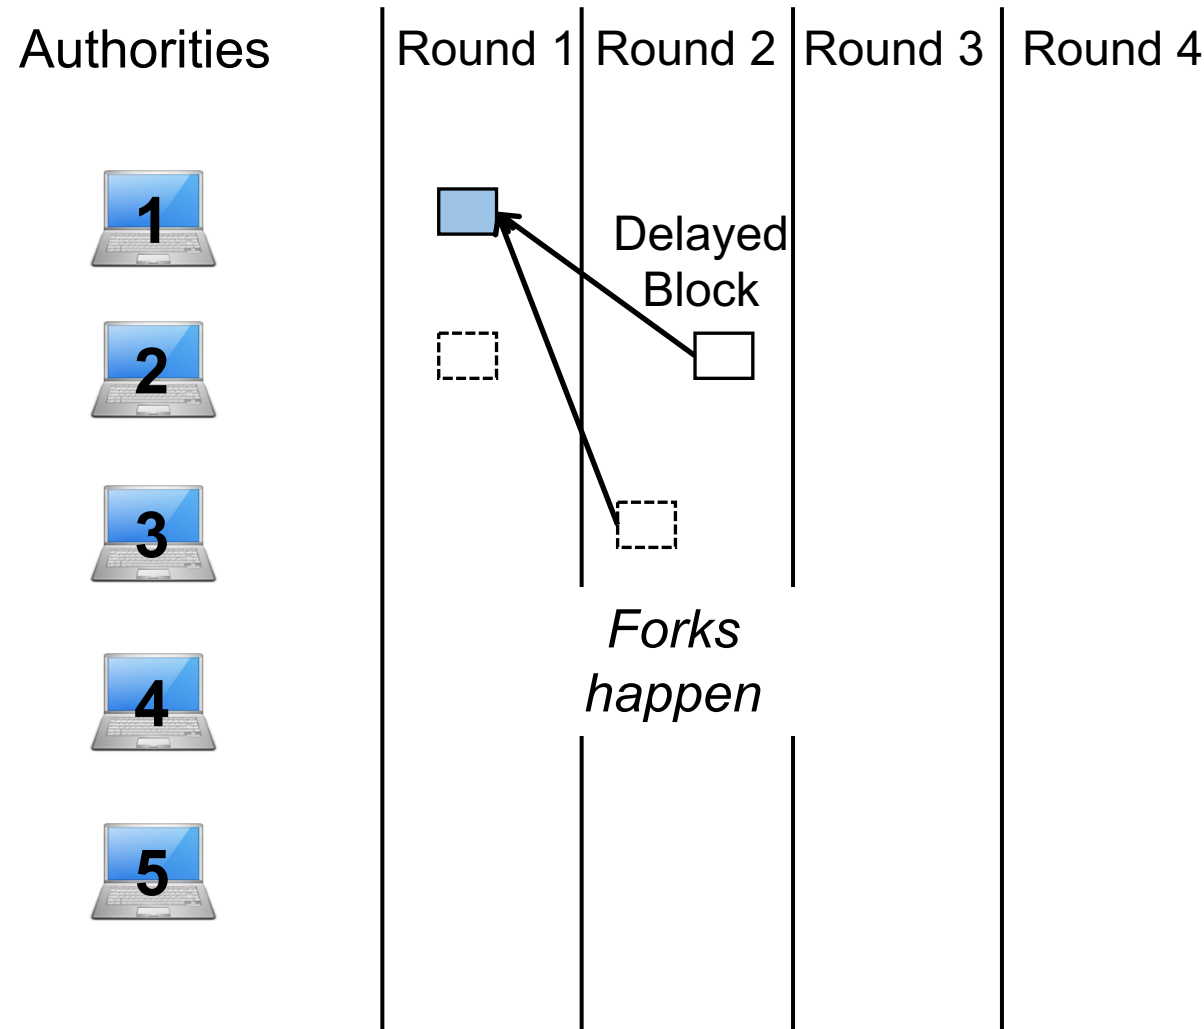
Block proposal

 by leader  Accepted

 by non-leader

*#non-leaders = up to $\left\lceil \frac{N}{2} \right\rceil - 2$

PoA Clique: An Illustration



- Round-based
 - (rotating) leader & round timeout Δ
- Multiple block proposals in a round
 - by a leader
 - by non-leaders* with random delays
- Single block acceptance in a round
 - Normal: a leader block

Block proposal



by leader



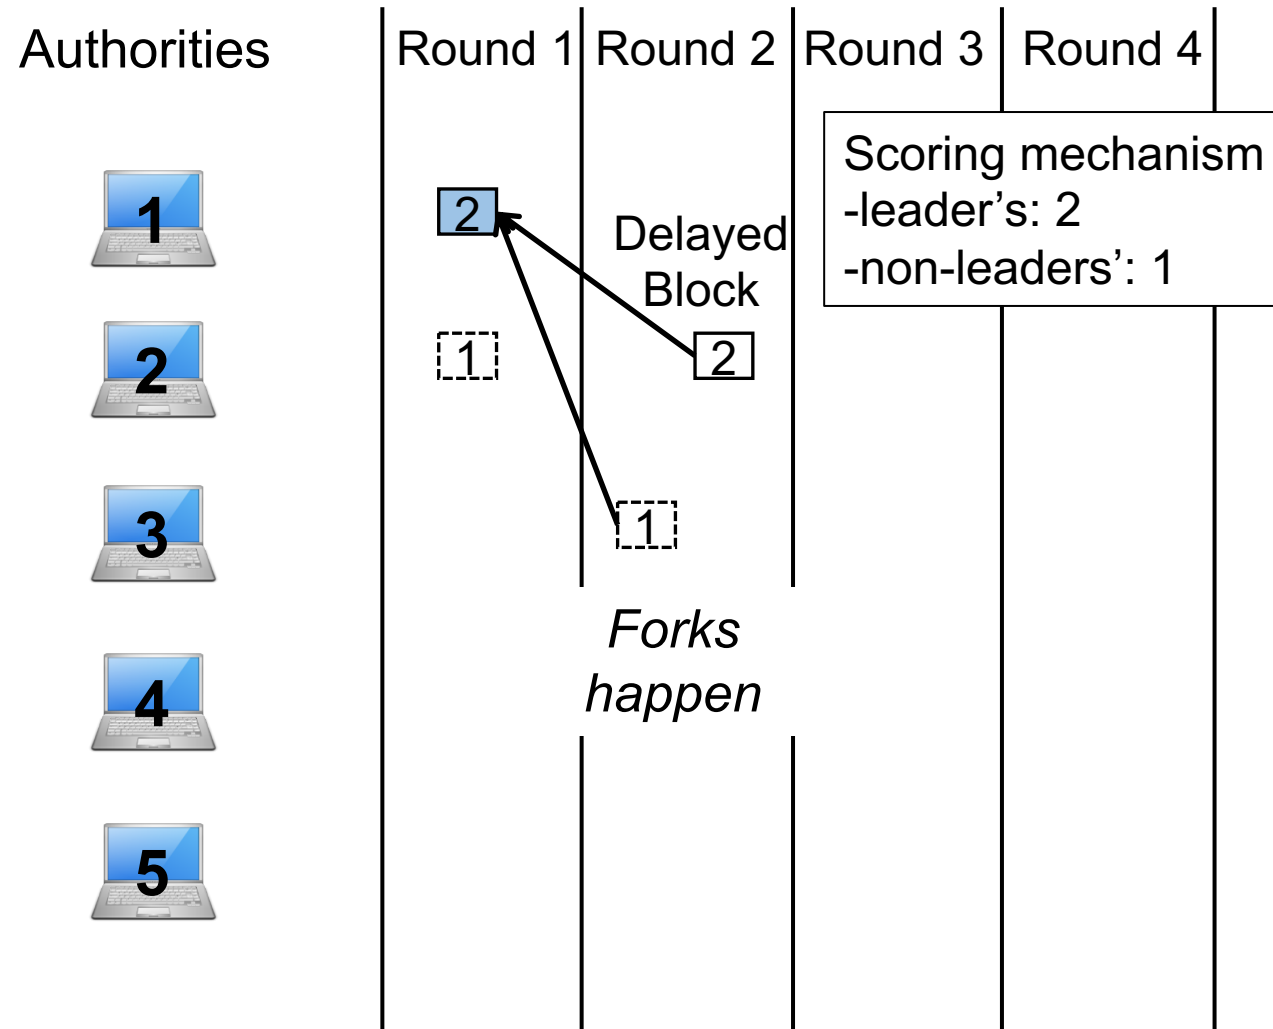
Accepted



by non-leader

*#non-leaders = up to $\left\lfloor \frac{N}{2} \right\rfloor - 2$

PoA Clique: An Illustration



- Round-based
 - (rotating) leader & round timeout Δ
- Multiple block proposals in a round
 - by a leader
 - by non-leaders* with random delays
- Single block acceptance in a round
 - Normal: a leader block
 - Forks: a highest scored chain

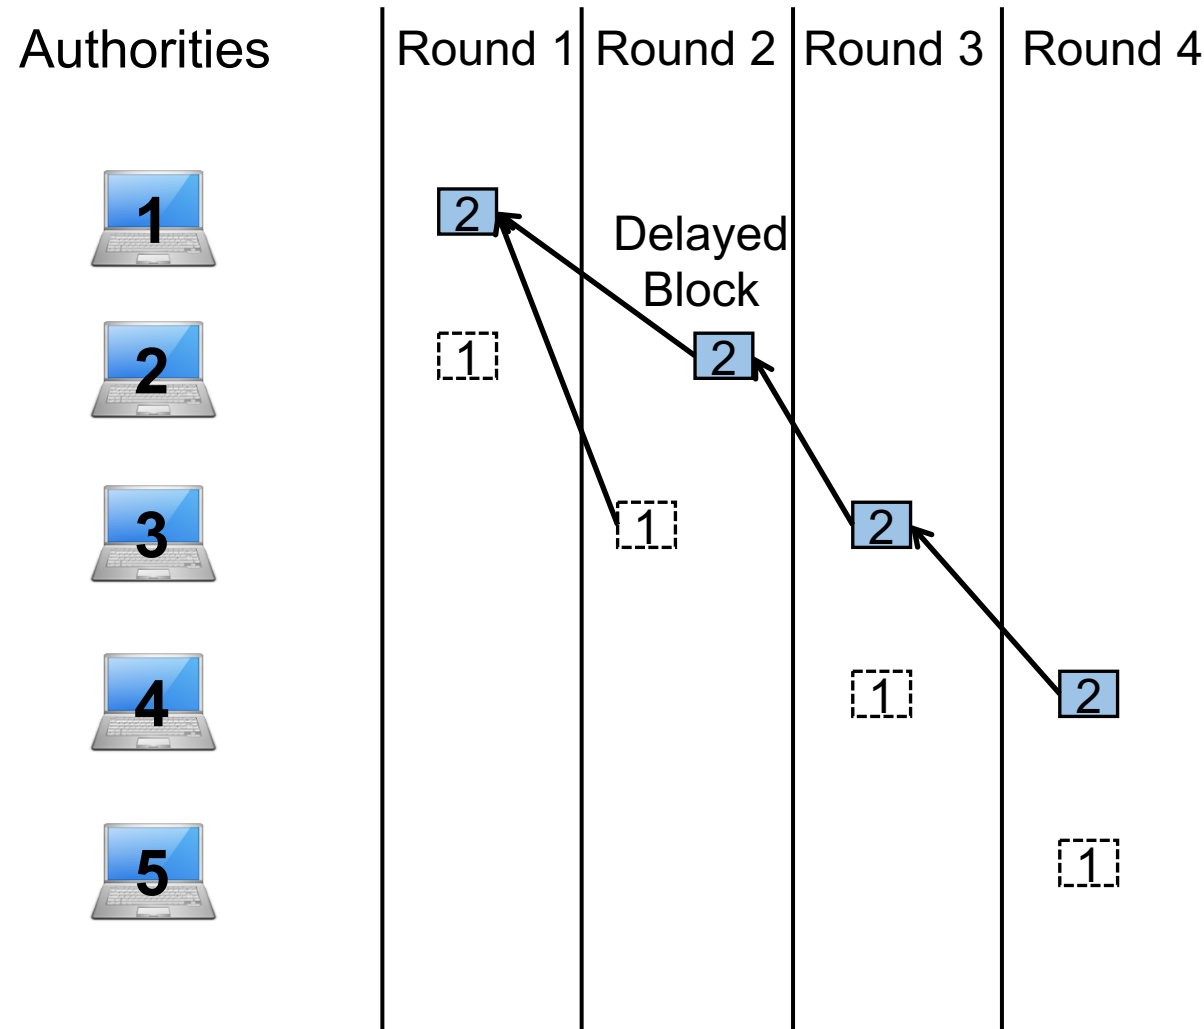
Block proposal

by leader
 Accepted

by non-leader

*#non-leaders = up to $\left\lfloor \frac{N}{2} \right\rfloor - 2$

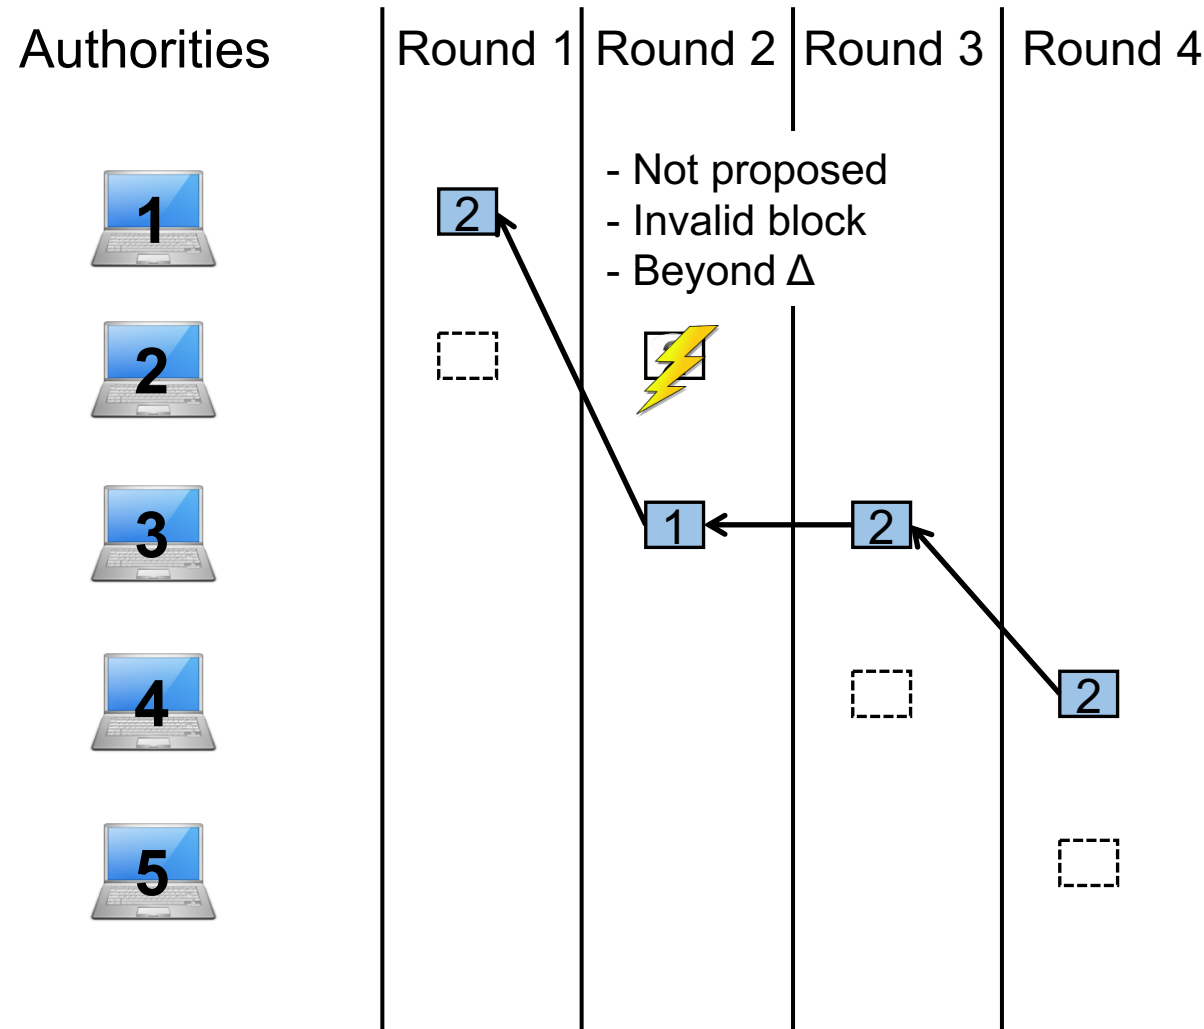
PoA Clique: An Illustration



- Round-based
 - (rotating) leader & round timeout Δ
- Multiple block proposals in a round
 - by a leader
 - by non-leaders* with random delays
- Single block acceptance in a round
 - Normal: a leader block
 - Forks: a highest scored chain

*#non-leaders = up to $\left\lceil \frac{N}{2} \right\rceil - 2$

PoA Clique: An Illustration



- Round-based
 - (rotating) leader & round timeout Δ
- Multiple block proposals in a round
 - by a leader
 - by non-leaders* with random delays
- Single block acceptance in a round
 - Normal: a leader block
 - Forks: a highest scored chain

Block proposal



by leader



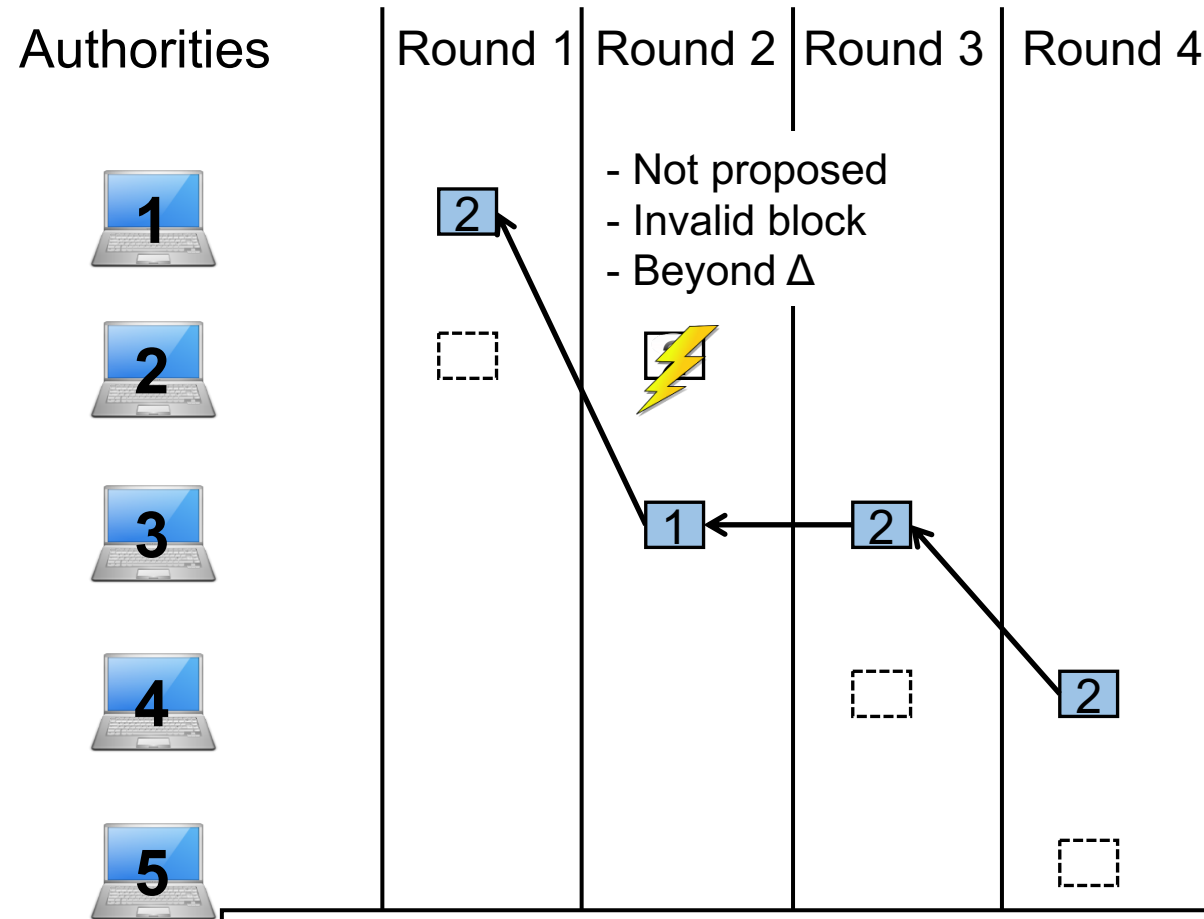
Accepted



by non-leader

*#non-leaders = up to $\left\lfloor \frac{N}{2} \right\rfloor - 2$

PoA Clique: Issues



- Round-based
 - (rotating) leader & round timeout Δ
- Multiple block proposals in a round
 - by a leader
 - by non-leaders* with random delays
- Single block acceptance in a round
 - Normal: a leader block
 - Forks: a highest scored chain

Multiple block proposals, but a single block acceptance

Block proposal



by leader




Accepted



by non-leader

*#non-leaders = up to $\left\lceil \frac{N}{2} \right\rceil - 2$

PoA Clique: Issues

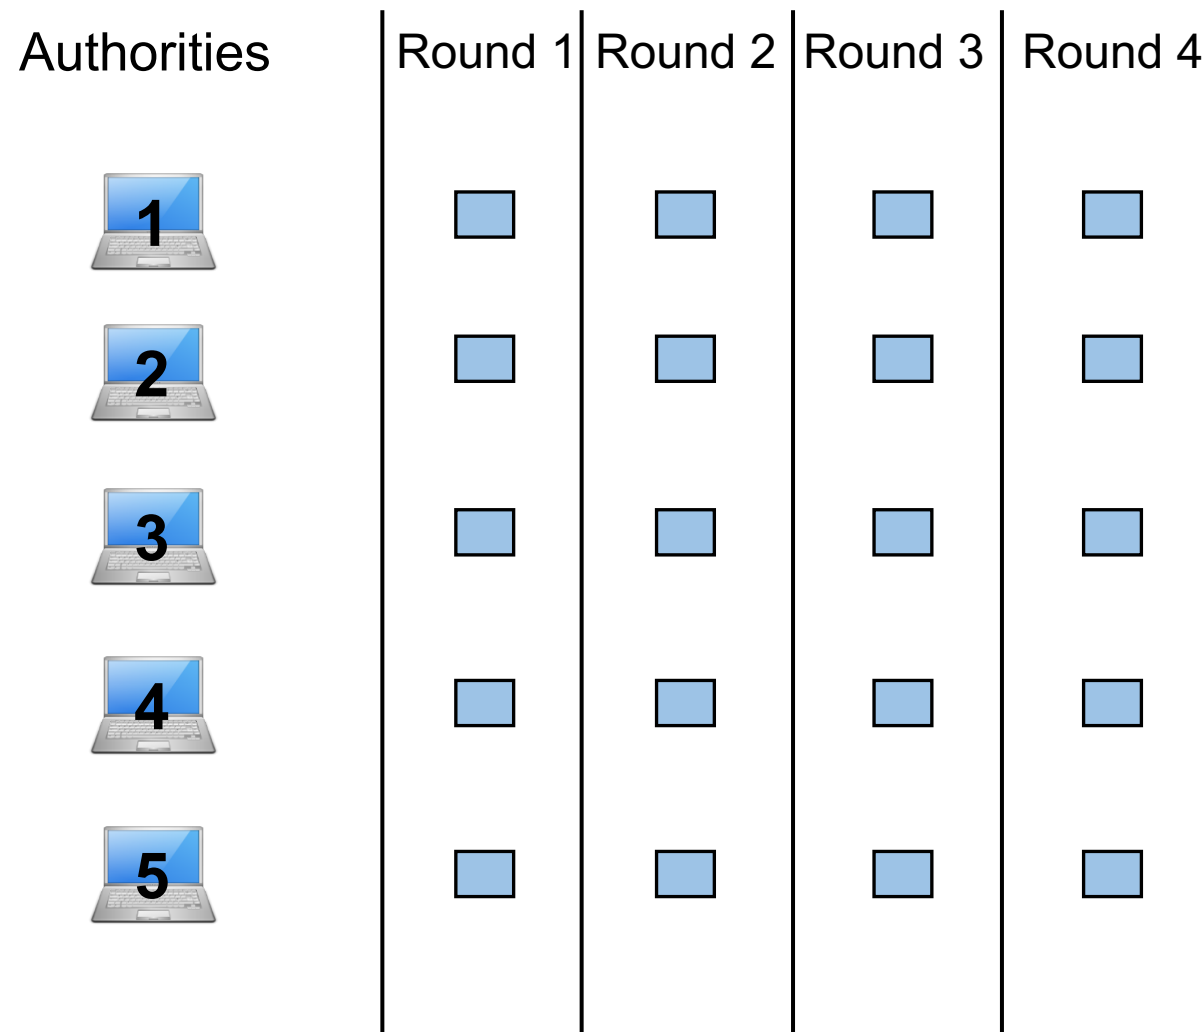


- Round-based
 - (rotating) leader & round timeout Δ
- Multiple block proposals in a round
 - by a leader
 - by non-leaders* with random delays
- Single block acceptance in a round
 - Normal: a leader block
 - Forks: a highest scored chain

Multiple block proposals, but a single block acceptance

Limited throughput & Wasted resources

Our Idea



Block proposal



by leader



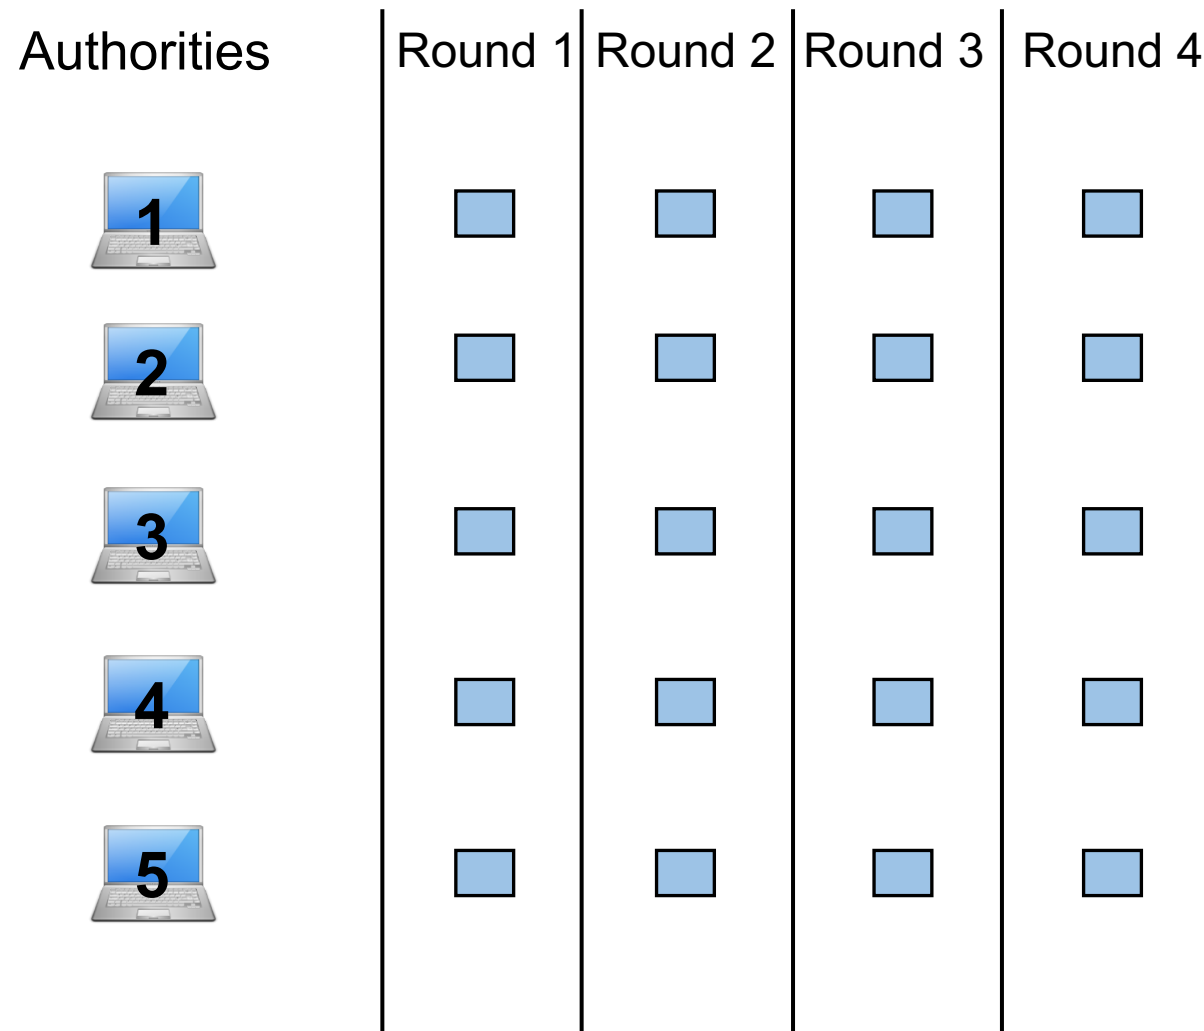
Accepted



by non-leader

- Round-based
 - leader-less & round timeout Δ
- Multiple block proposals in a round
 - By all authorities
- Multiple block acceptance in a round
 - Normal: (consistent) all proposed blocks
 - Forks: (consistent) total order

Our Idea



Block proposal



by leader



Accepted




























by non-leader

- Round-based
 - leader-less & round timeout Δ
- Multiple block proposals in a round
 - By all authorities
- Multiple block acceptance in a round
 - Normal: (consistent) all proposed blocks
 - Forks: (consistent) total order

↓
Challenges
























Challenge (1) Consistent view across nodes

Authority 1's view

Authorities	Round 1	Round 2	Round 3	Round 4
				
				
Fork attack 				
				
				

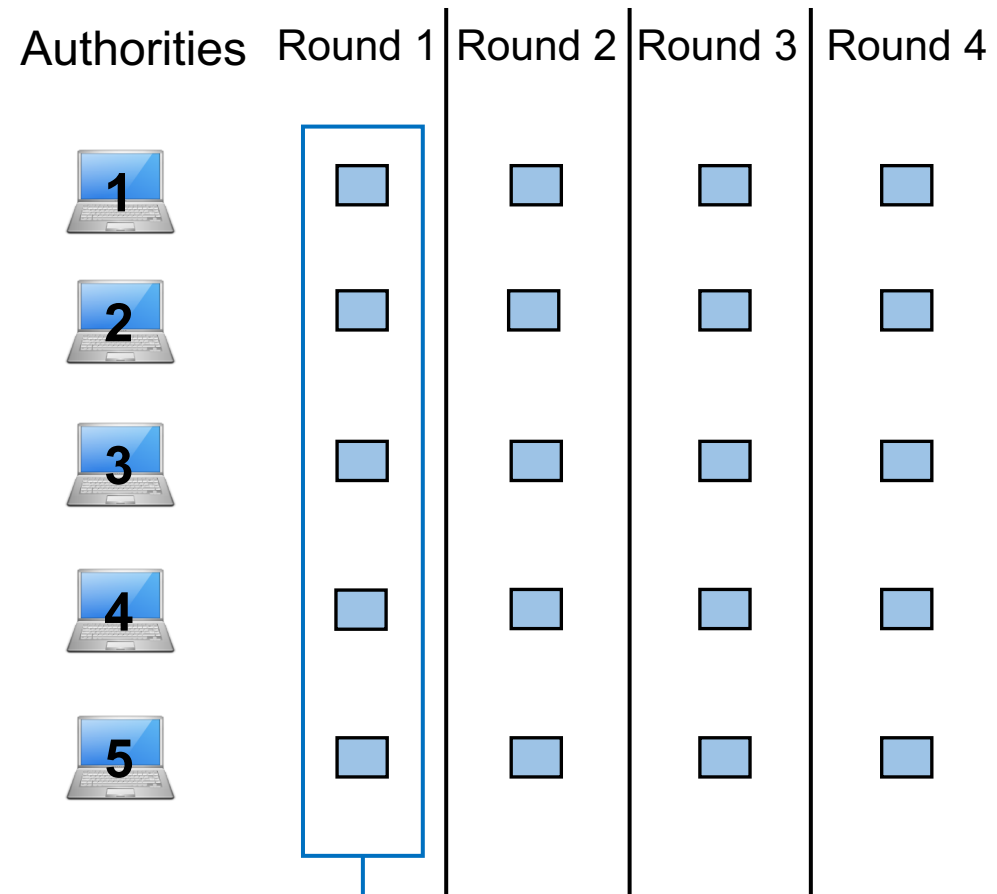


Authority 2's view

Authorities	Round 1	Round 2	Round 3	Round 4
				
				
				
		<div>Delayed/ Dropped</div>		
				

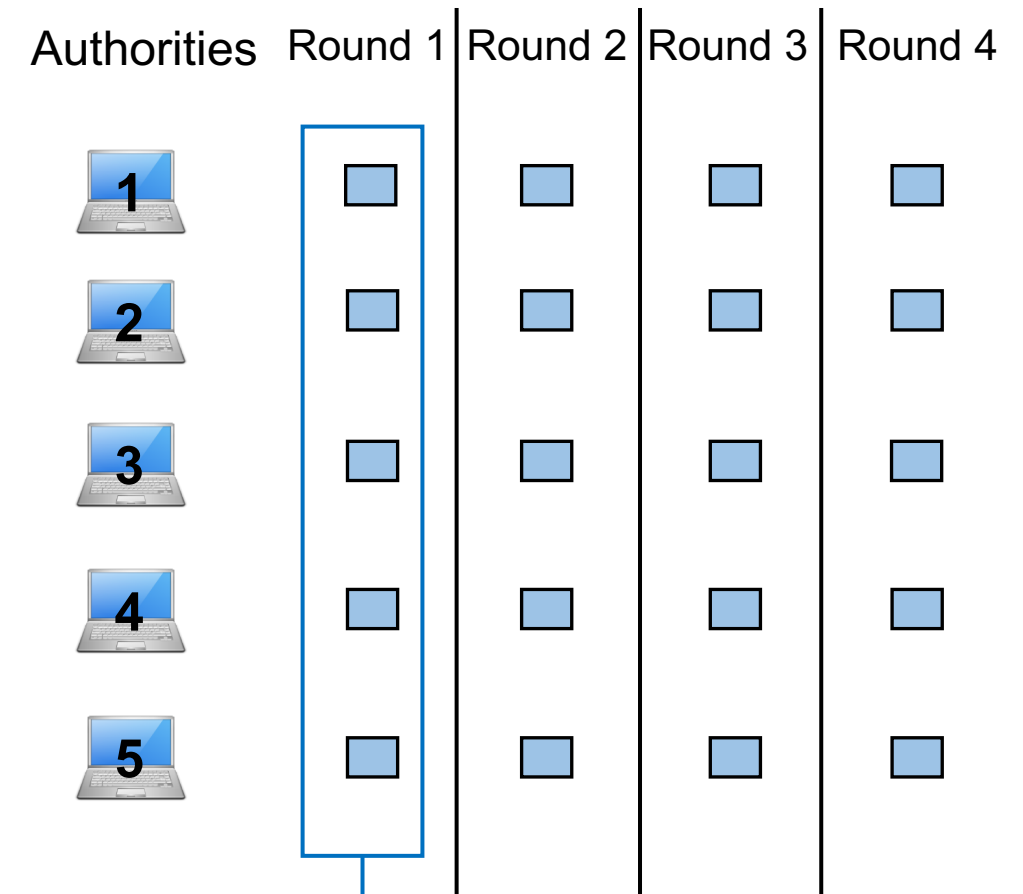
Challenge (2) Consistent total ordering

Authority 1's view



Authority 1's
total order

Authority 2's view




Authority 2's
total order

=


?
=

DAG-based BFT Consensus


DAG Mempool



- Consistent blocks across nodes
- Reliable block distribution
- Parallel block proposals & Acceptance

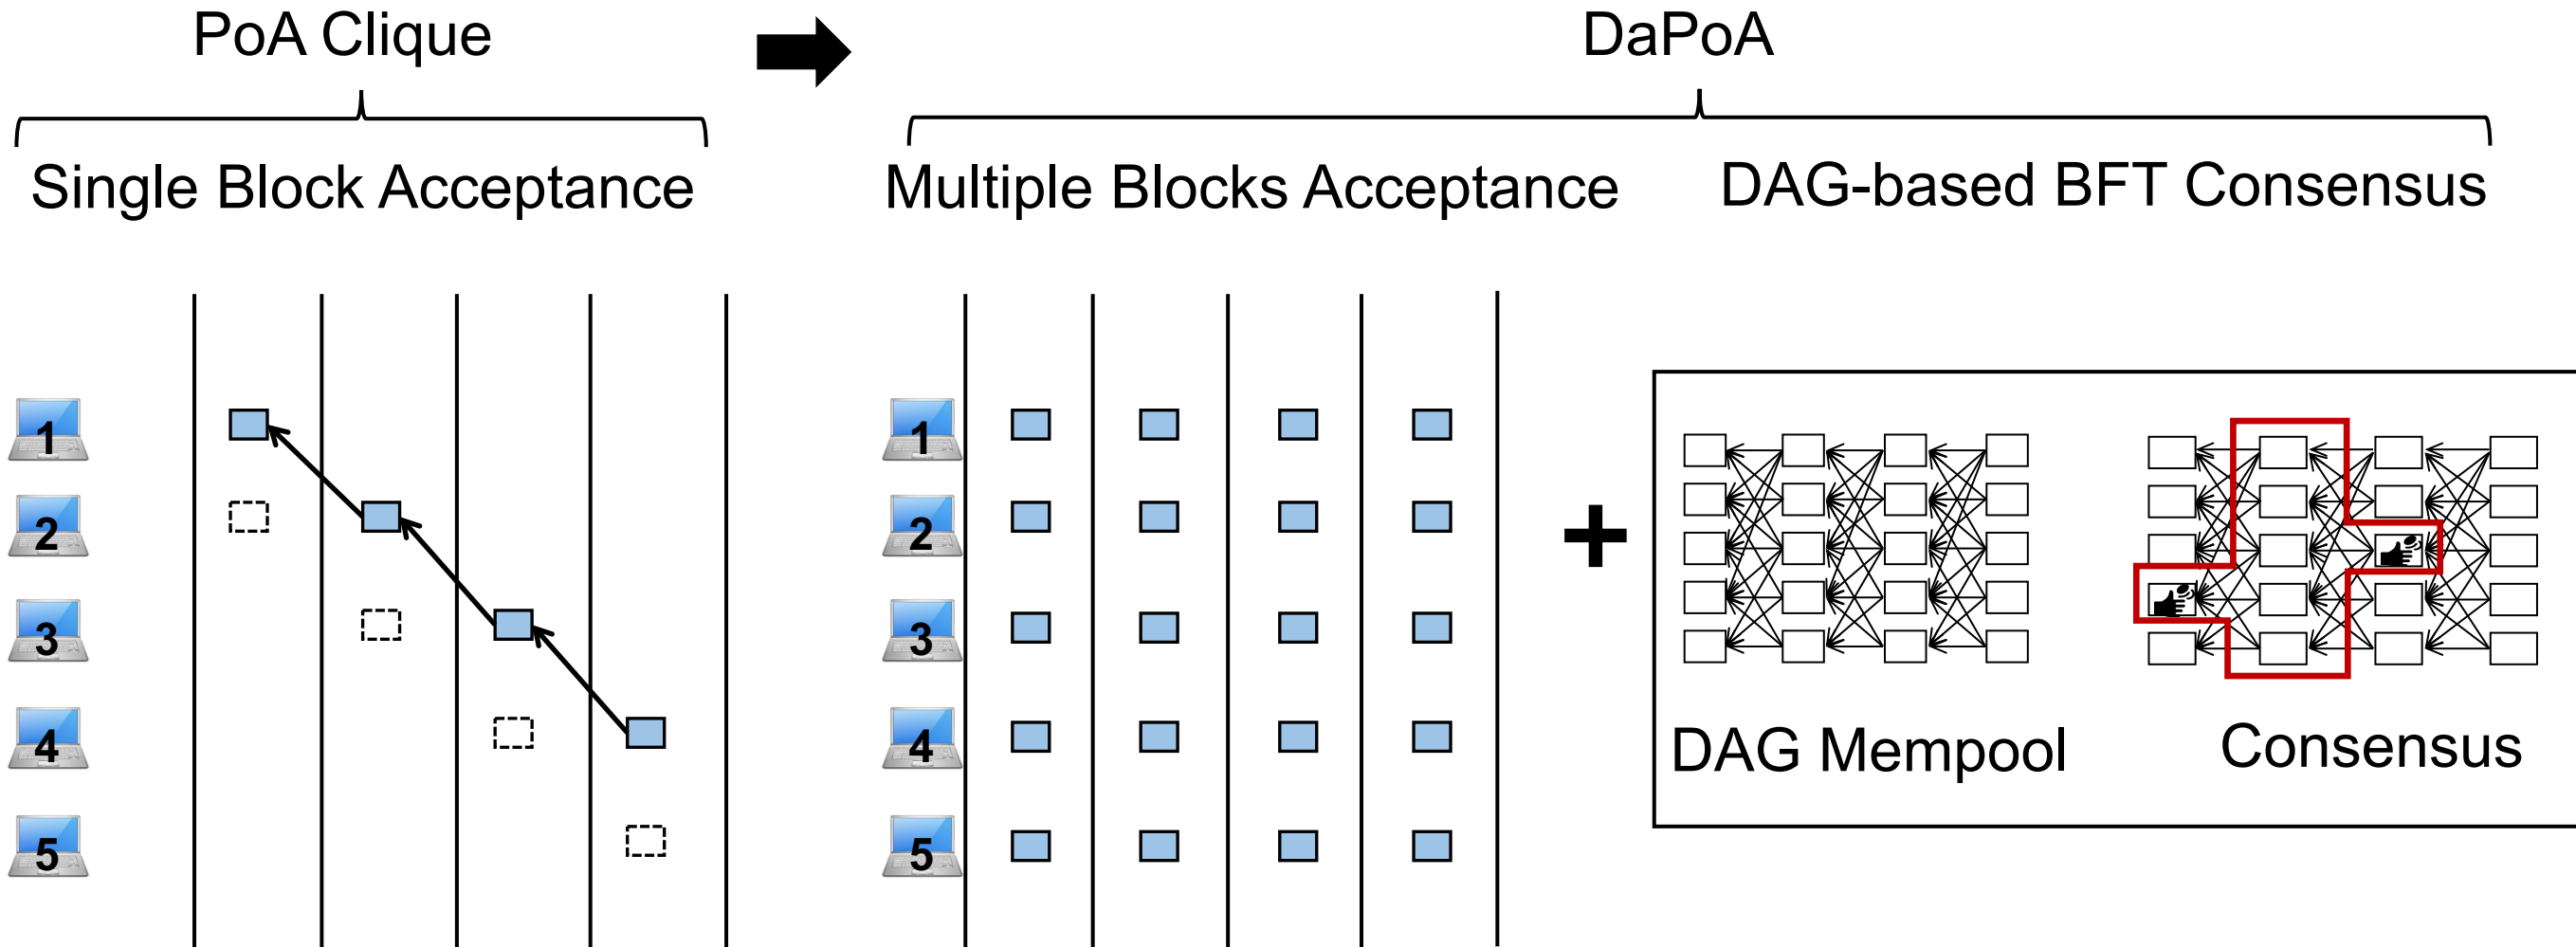


Consensus



- Consistent total order
- Local & deterministic
- (shared) randomness

Contributions





achieves 2.47x higher throughput, 5.76x lower latency than Clique


Assumptions

- A consortium blockchain
 - Known Identities
 - Members operates DaPoA node
 - Clients submit Txs to a trusted node
 - f Byzantine members out of $3f+1$ members
- Partially asynchronous network
 - unknown time bound Δ
- Crypto cannot be subverted

DaPoA: Operations

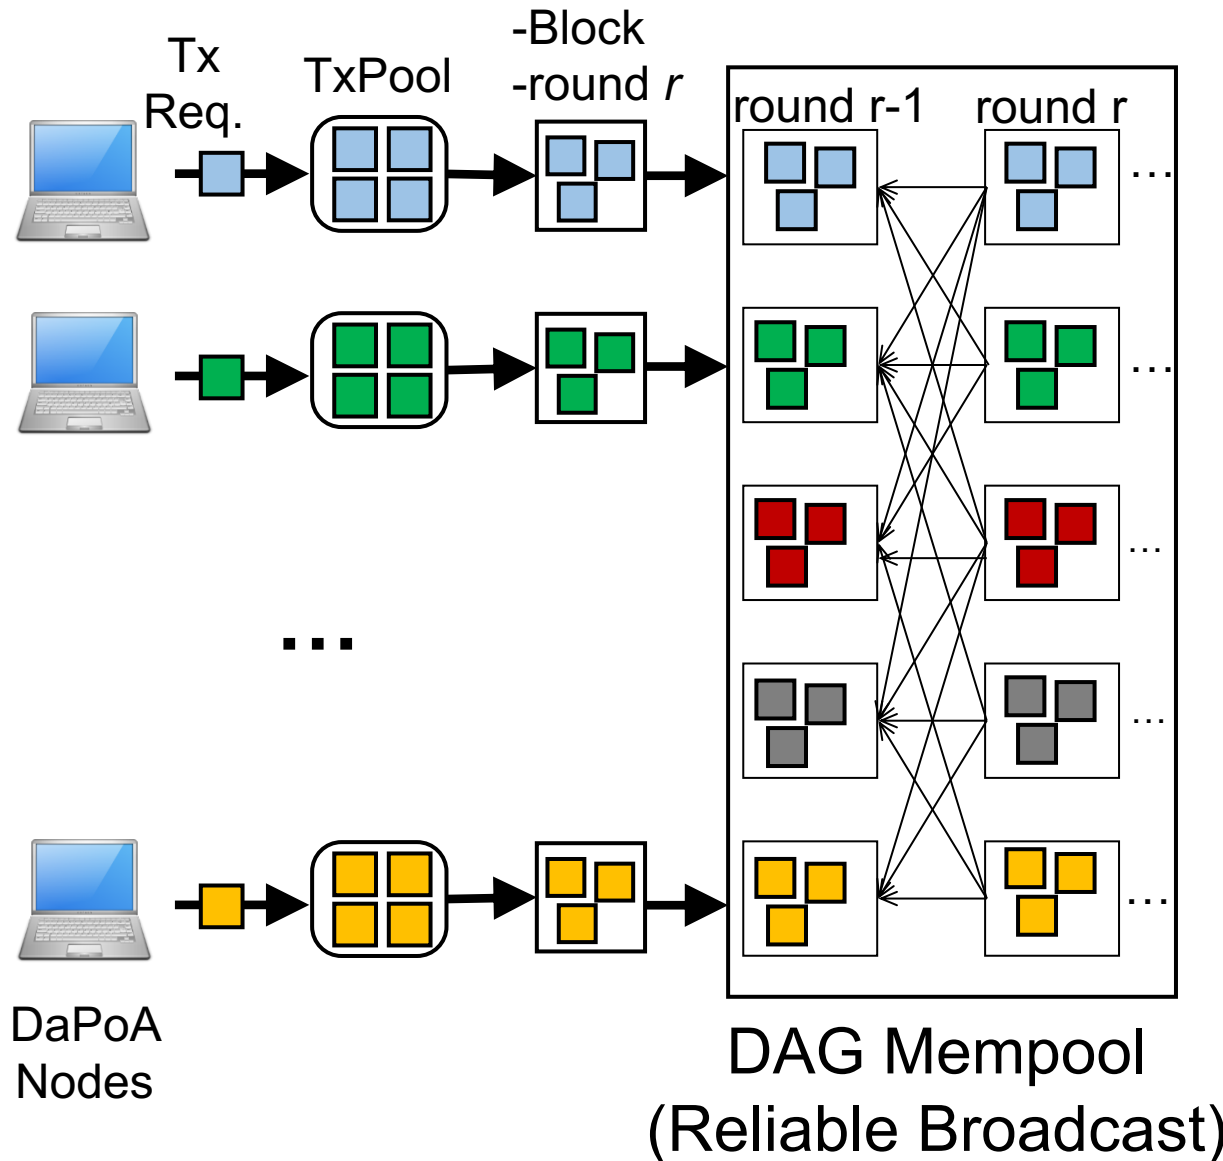


...



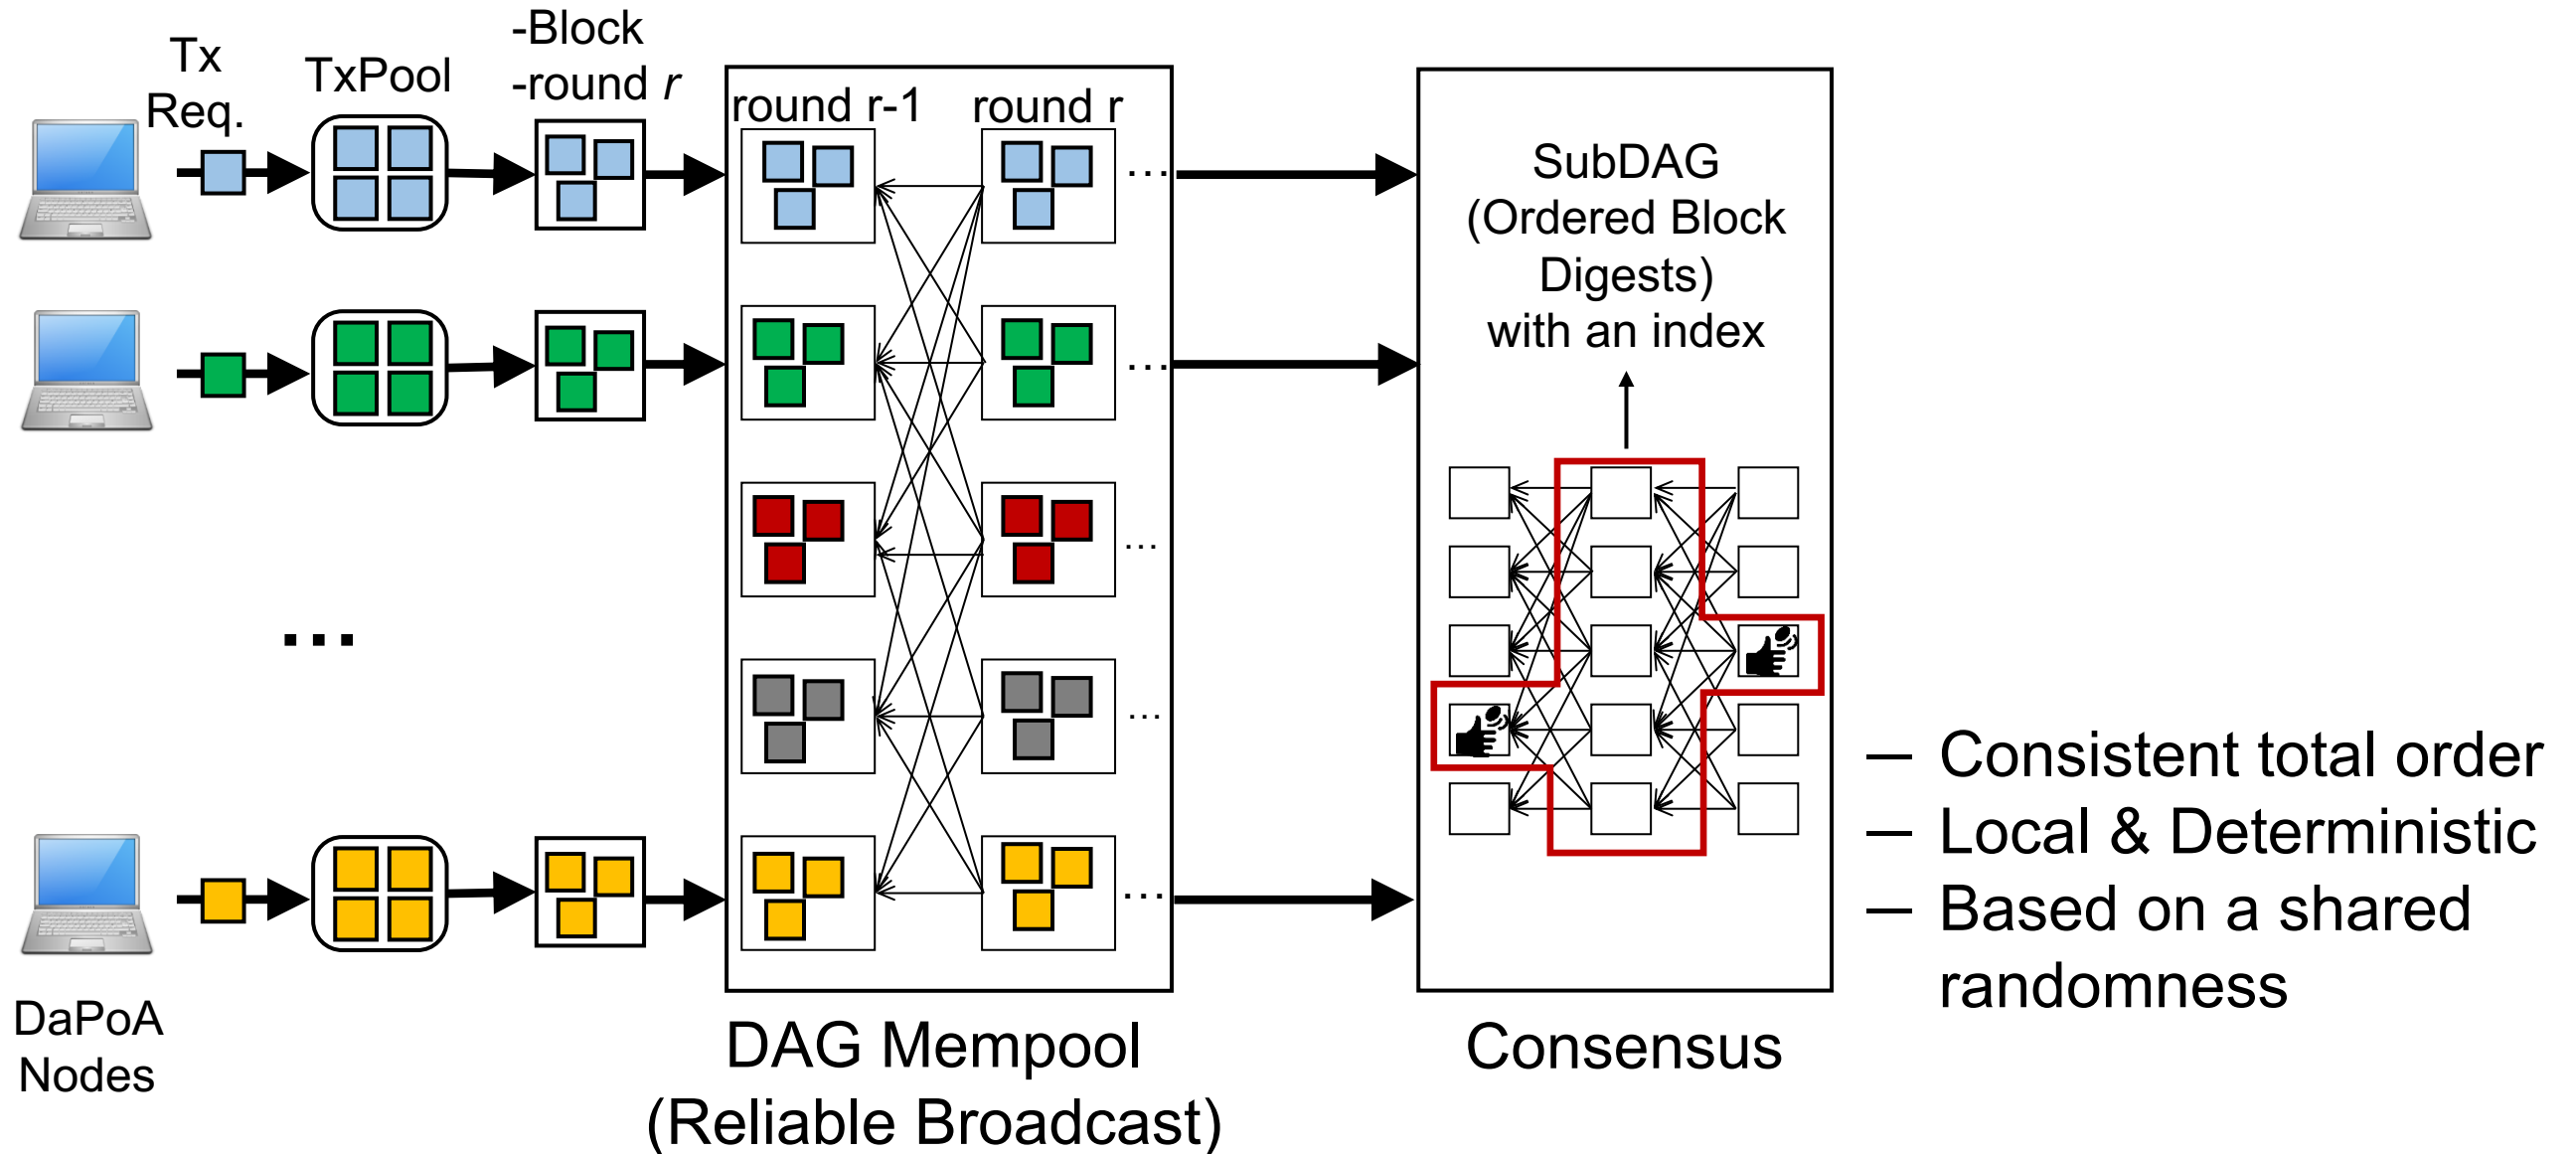
DaPoA
Nodes

DaPoA: Operations

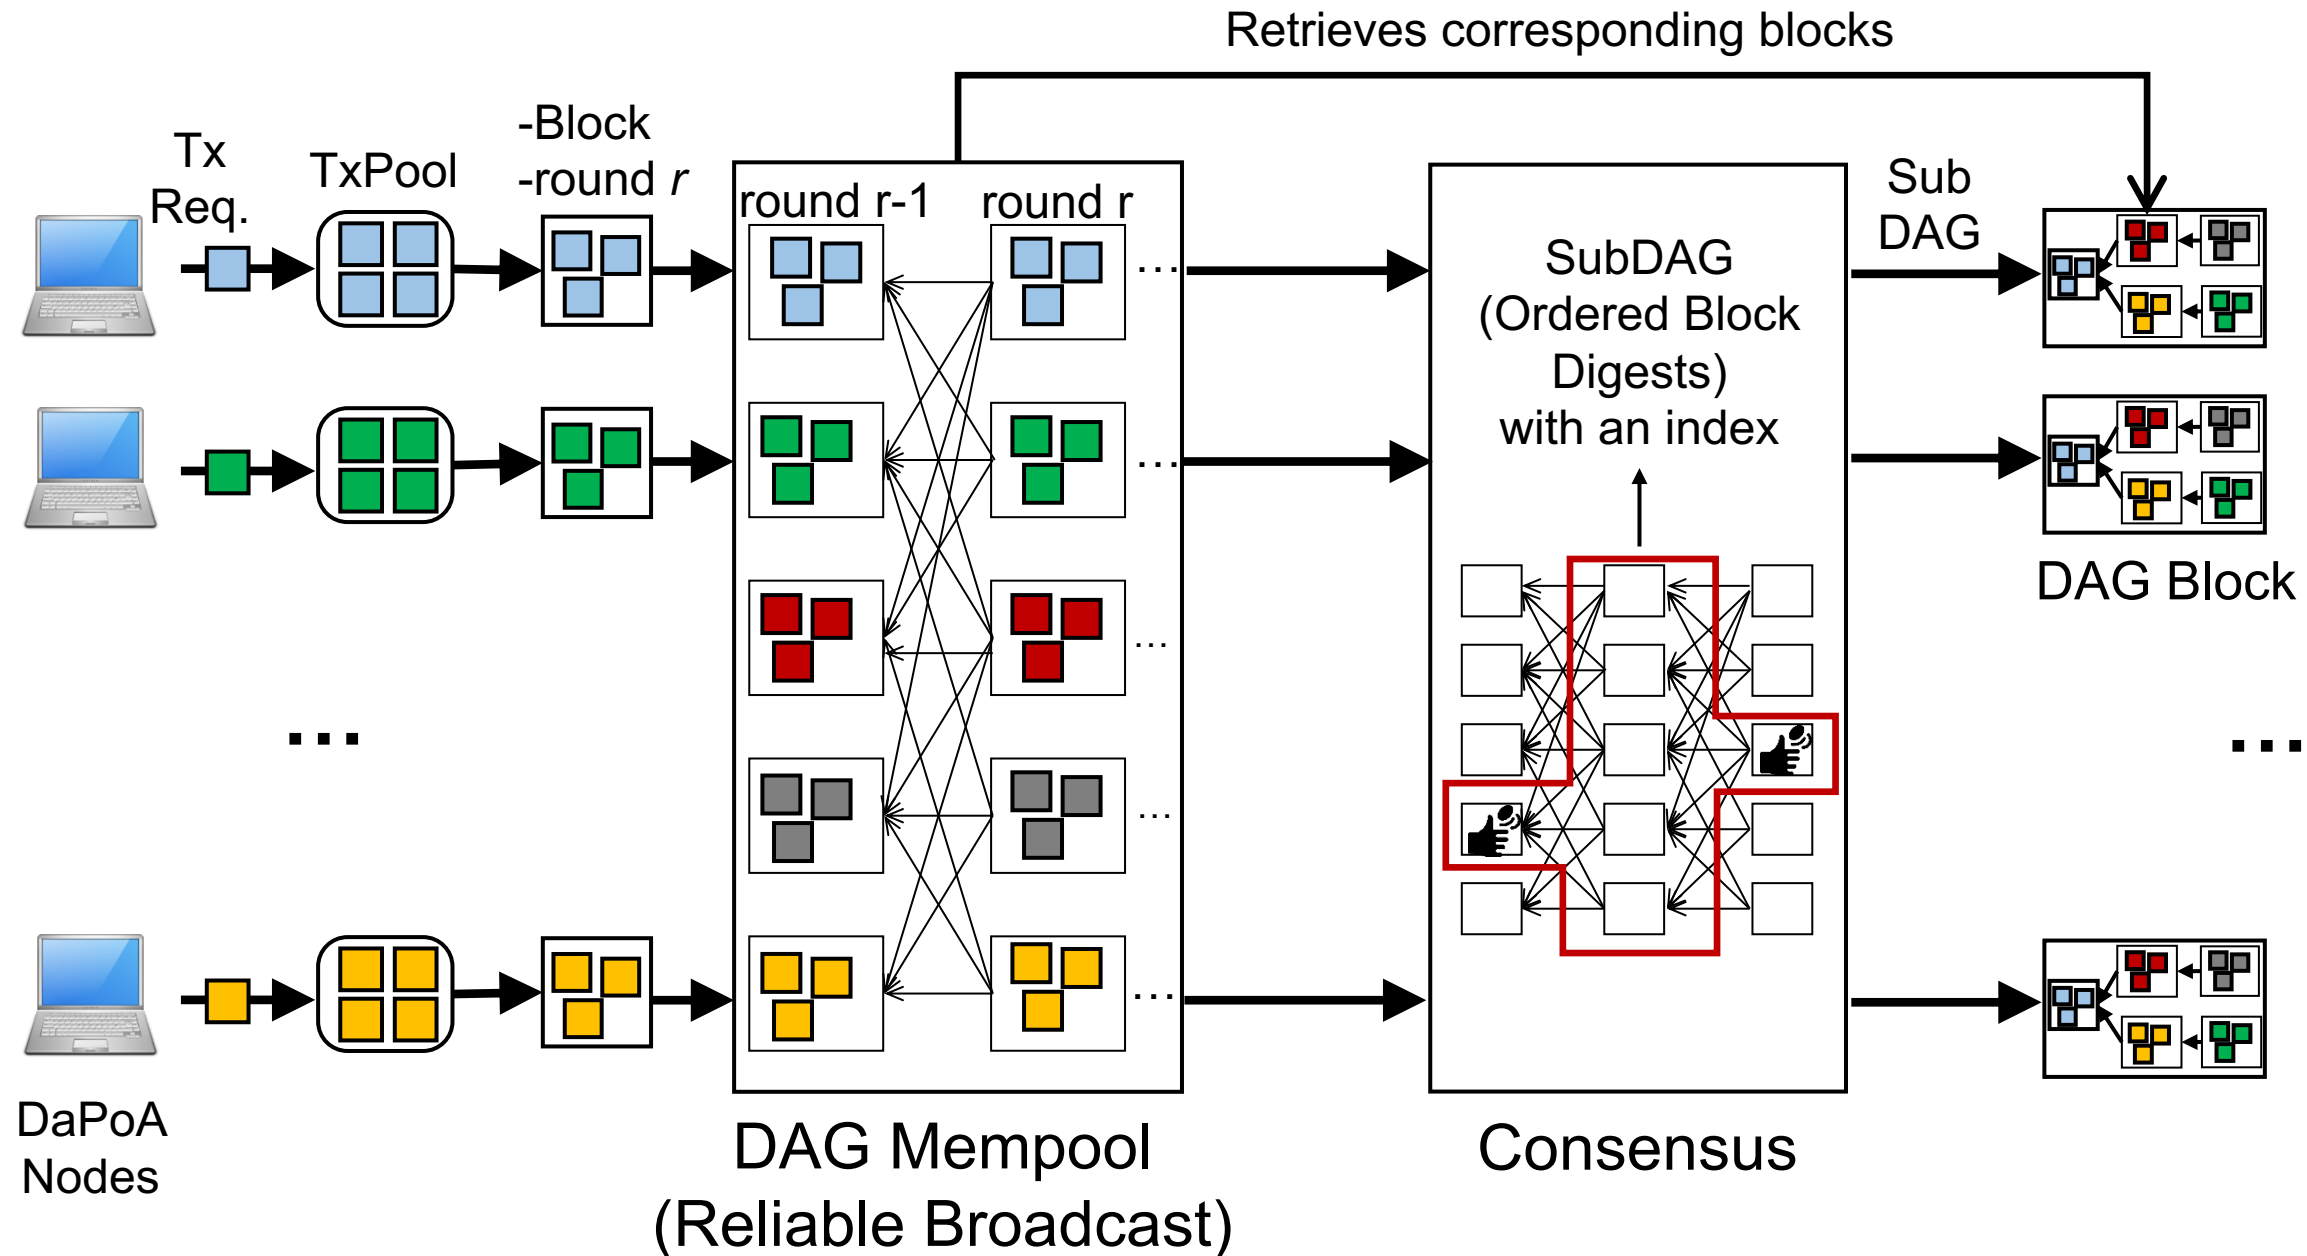


- Reliable block replication across nodes
- Consistent block views
- Multiple block proposals & acceptance

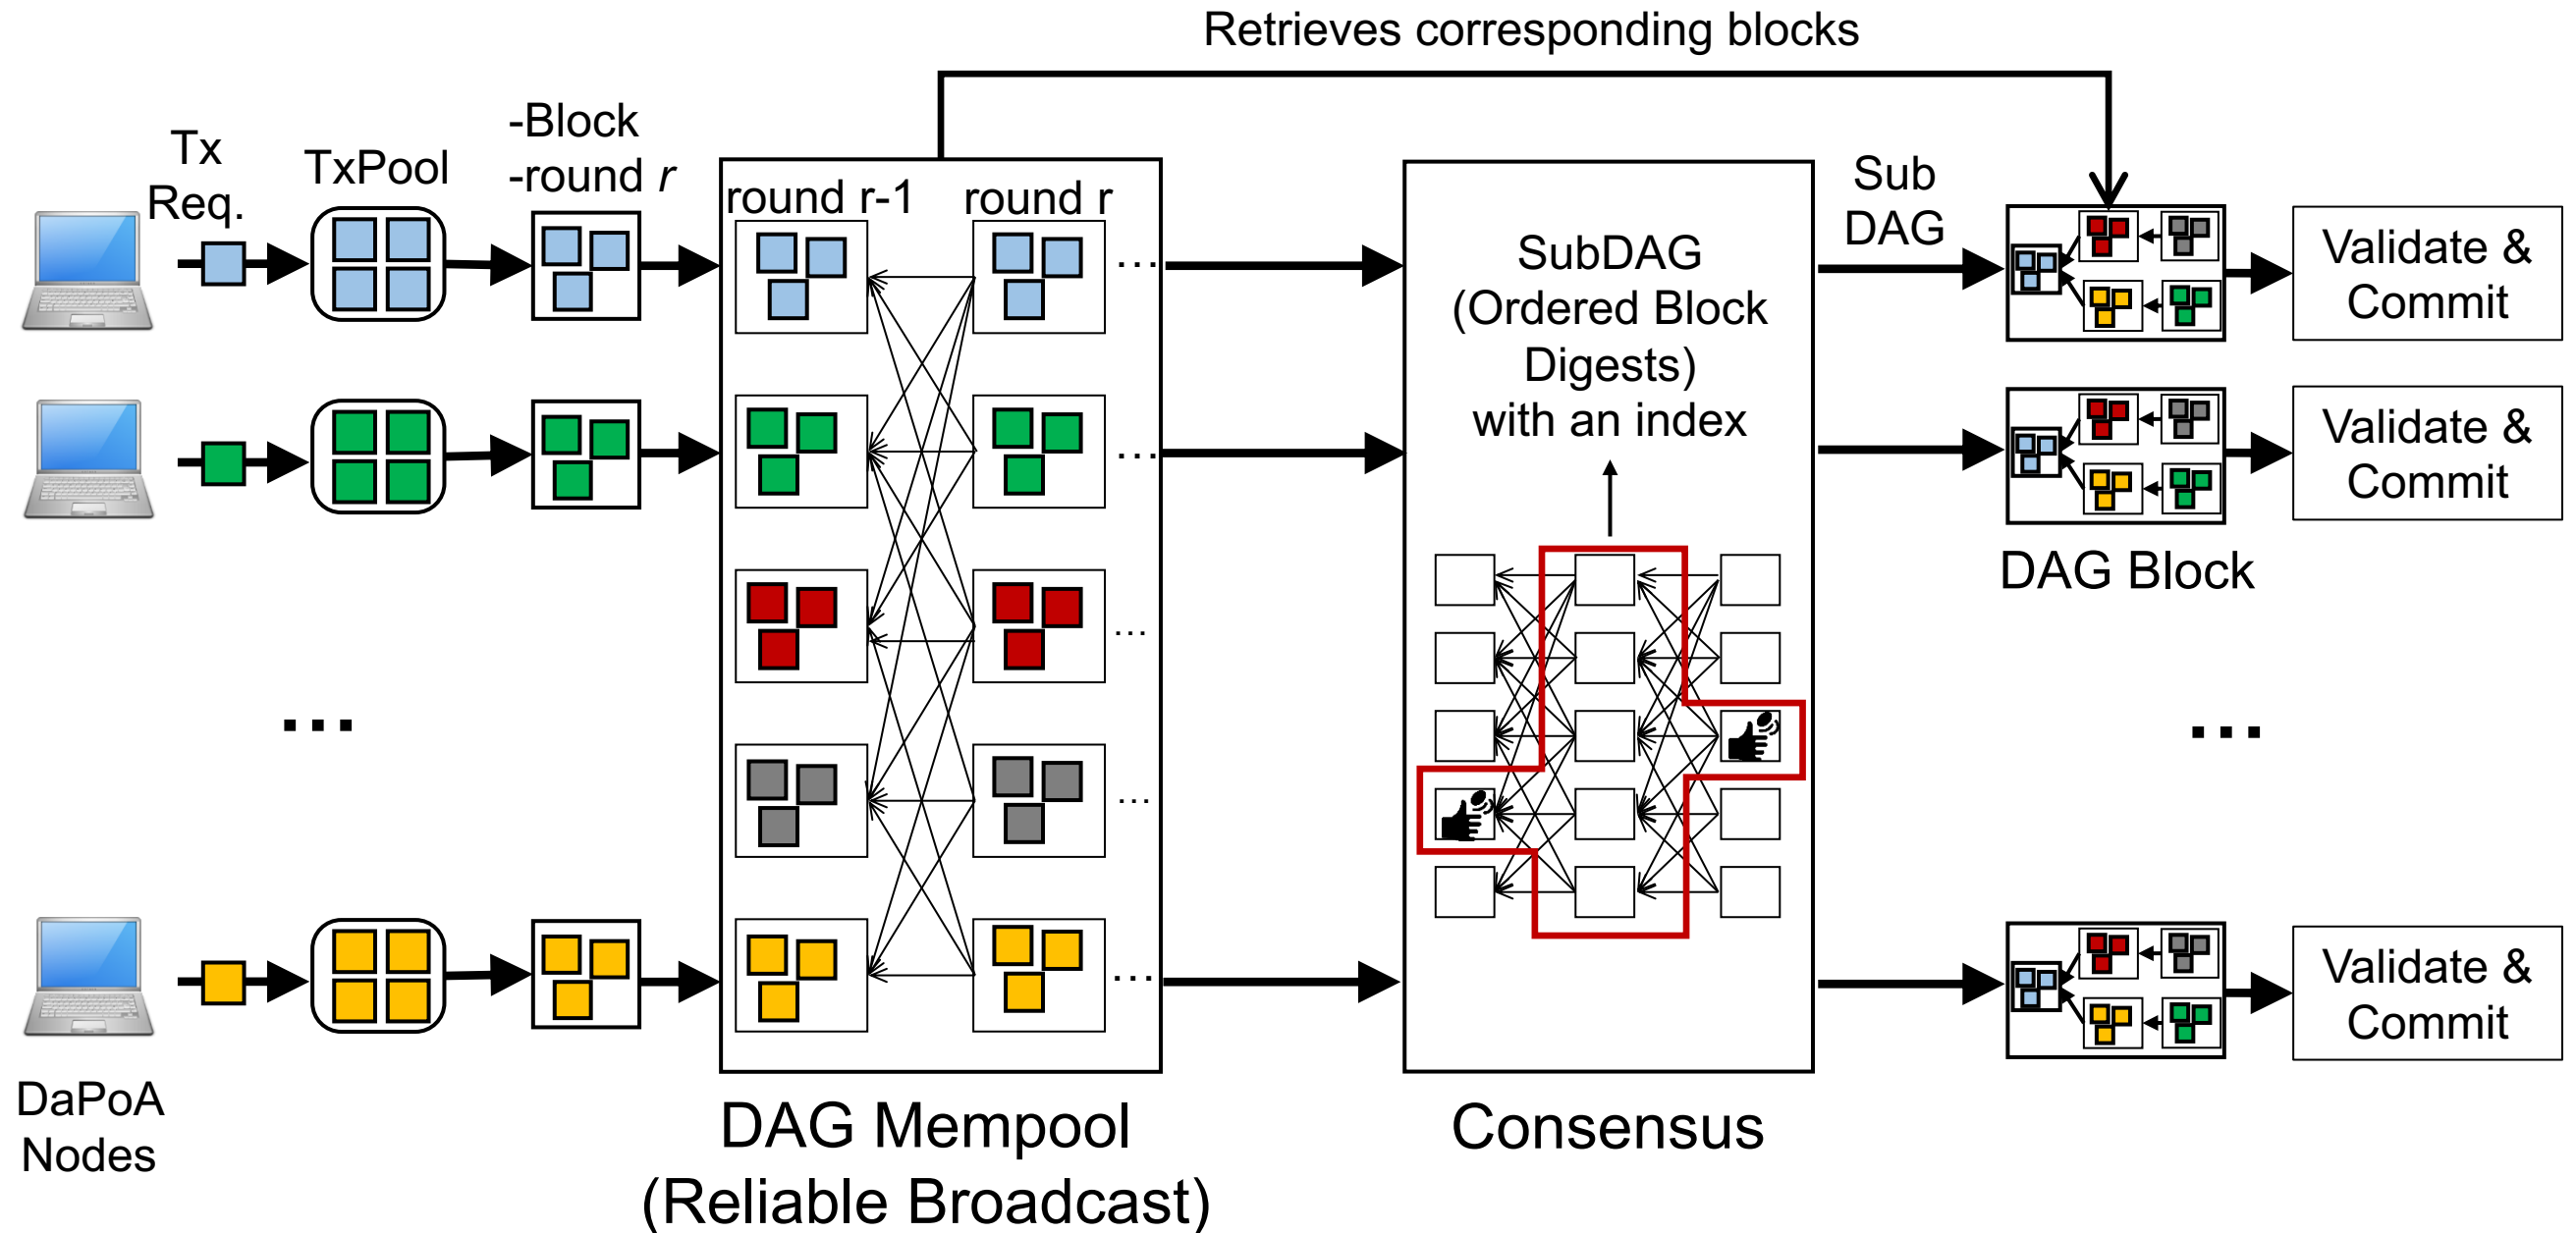
DaPoA: Operations



DaPoA: Operations




DaPoA: Operations



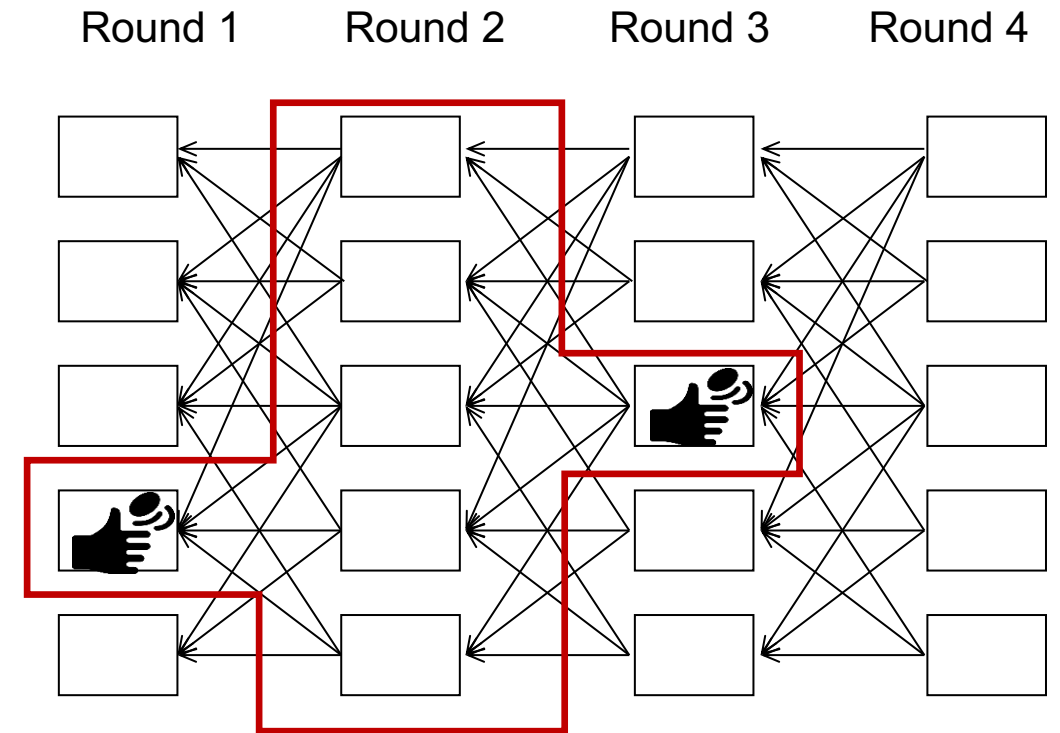
Security Analysis

Accept a proposal with $2f+1$ votes from the next round

Fork Attack



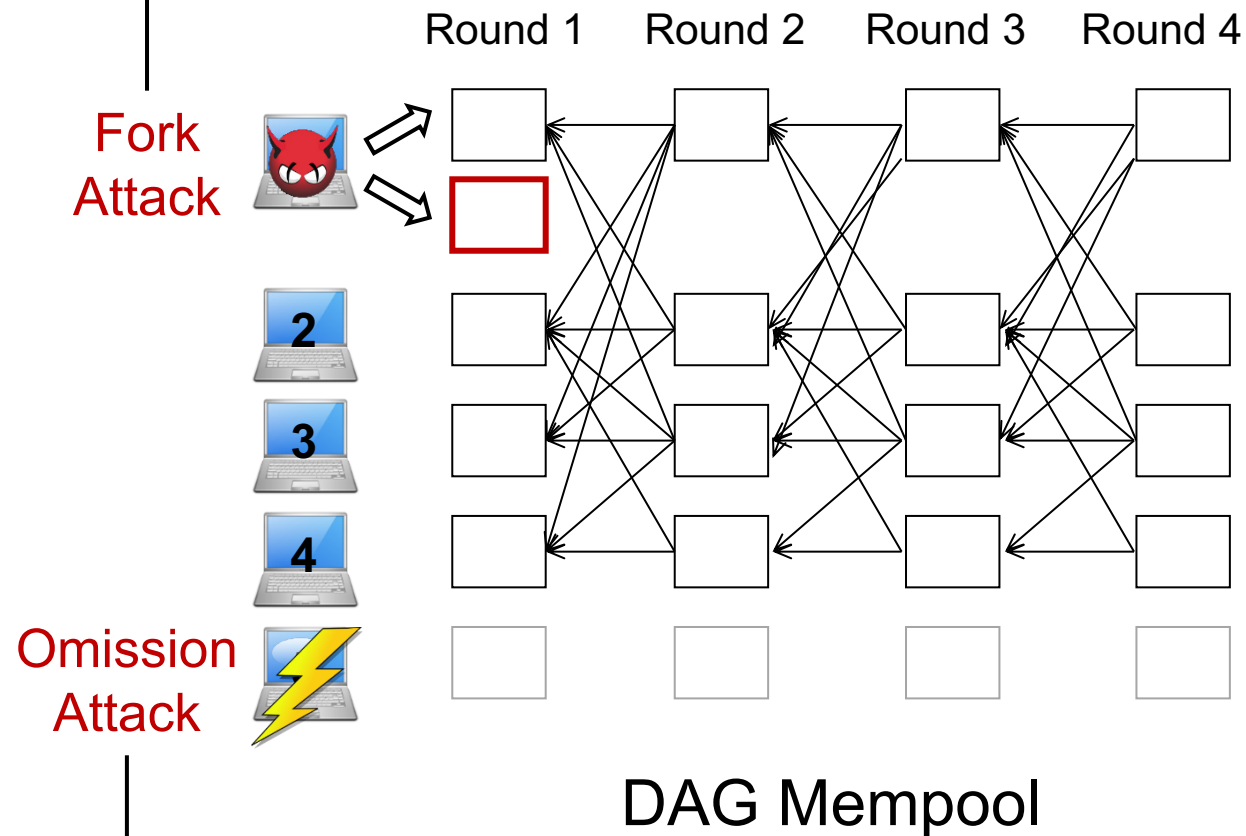
DAG Mempool



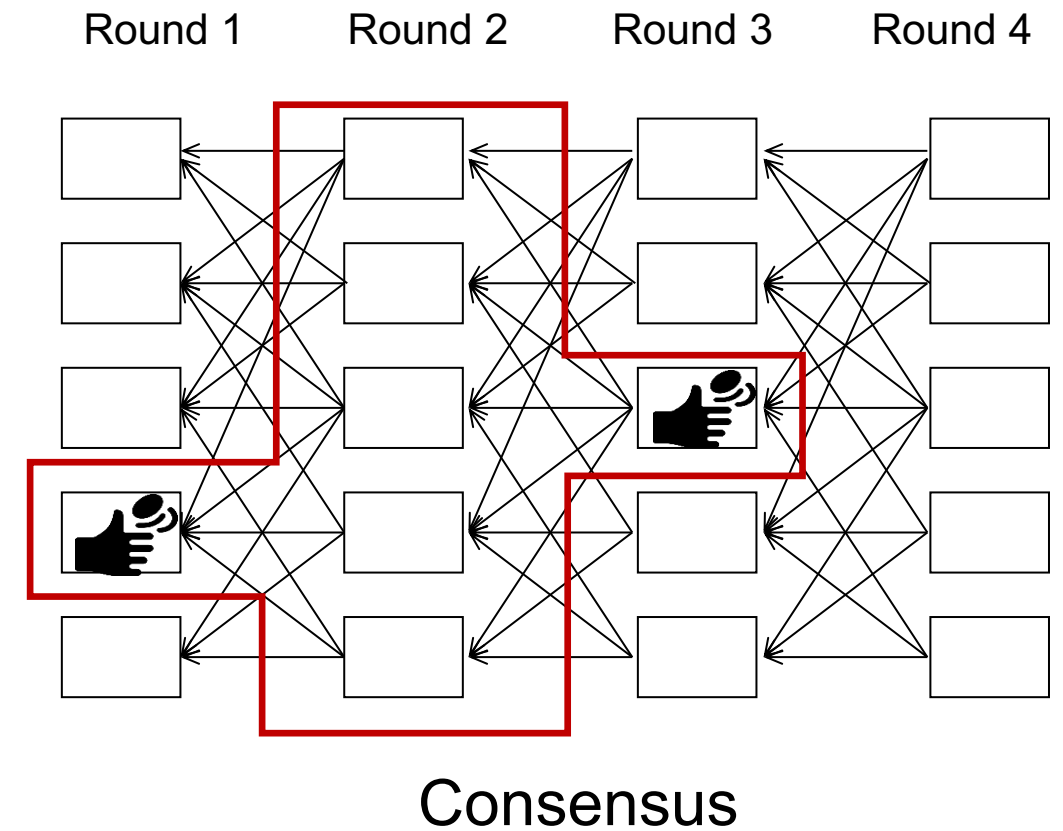
Consensus

Security Analysis

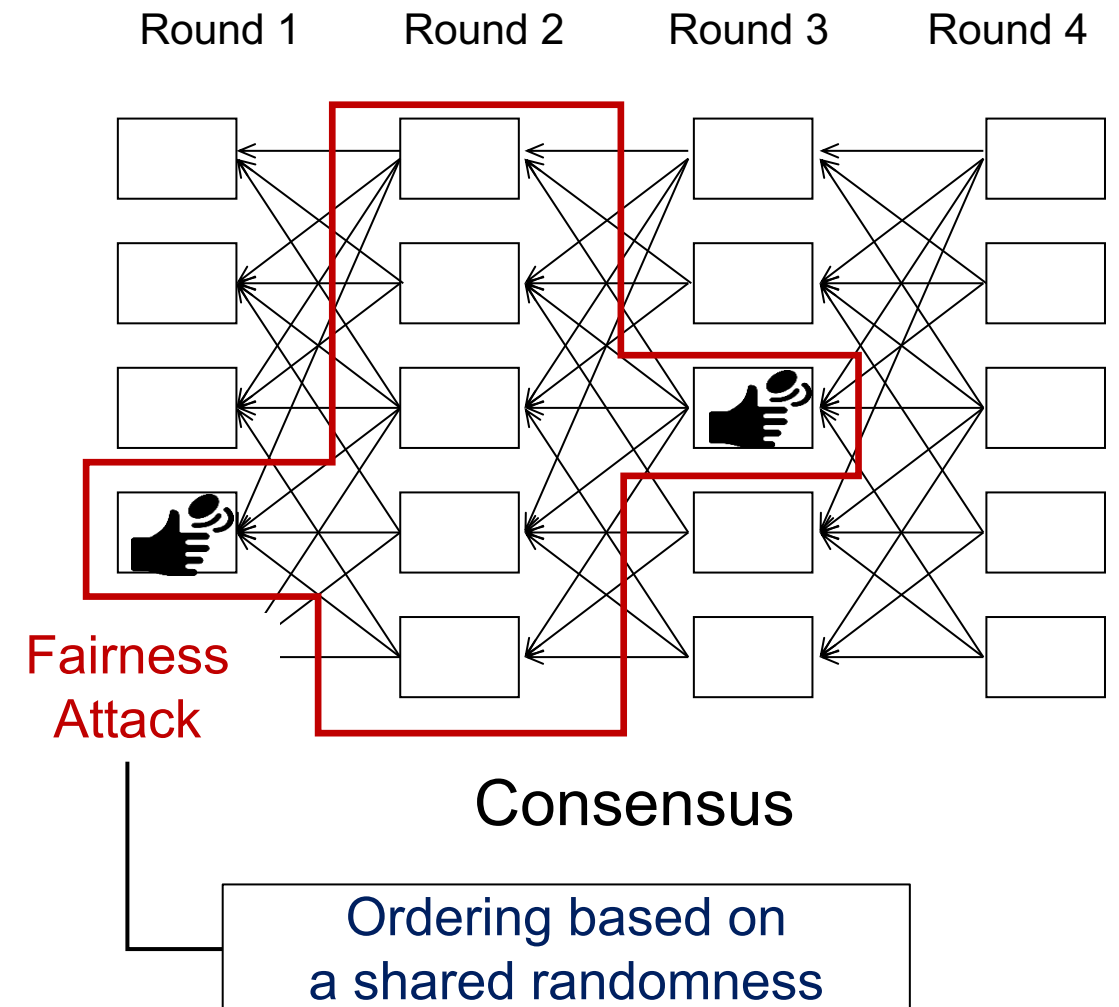
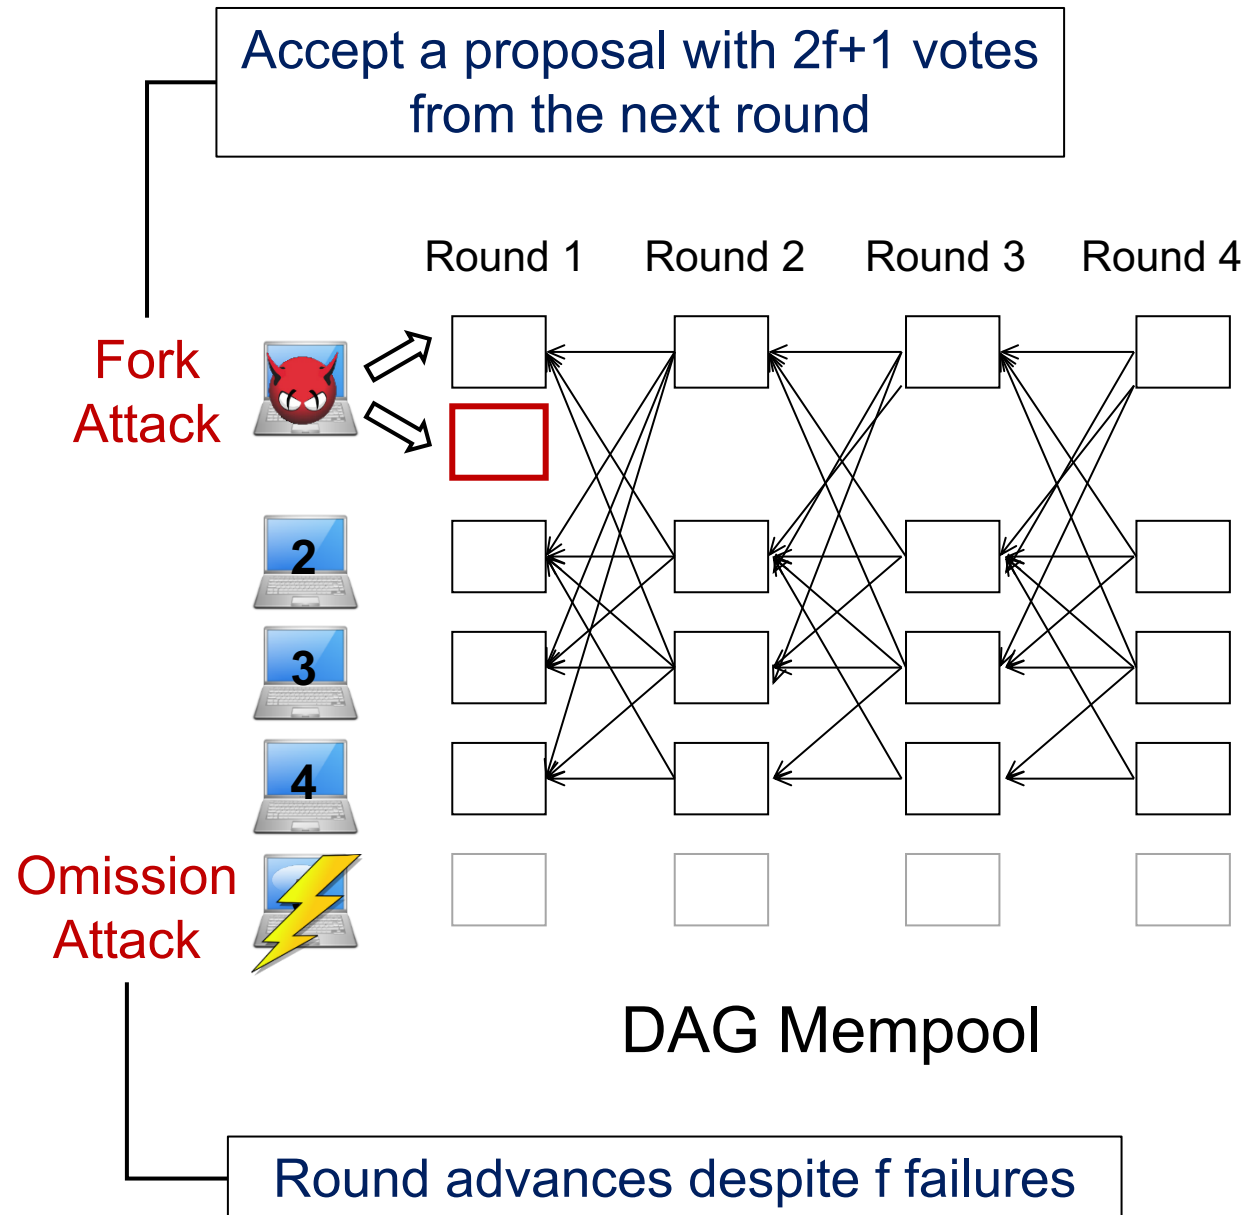
Accept a proposal with $2f+1$ votes from the next round



Round advances despite f failures



Security Analysis




Evaluation: vs. Clique

- Implement DaPoA based on Ethereum Geth*
- Narwhal/Tusk** for DAG BFT
- Benchmark using Hyperledger Caliper
- Hardware: AMD Ryzen 3990X CPU, 256 GB RAM

Parameters

- Simple payment contract
- TX confirmation blocks: 2
- Block Period: 1
- Gaslimit: 10^9
- DAG BFT (Batch Size: 2MB, 1 worker, 1 primary)
- Send rate: 700





Evaluation: vs. Clique

- Implement DaPoA based on Ethereum Geth*
- Narwhal/Tusk** for DAG BFT
- Benchmark using Hyperledger Caliper
- Hardware: AMD Ryzen 3990X CPU, 256 GB RAM

Parameters

- Simple payment contract
- TX confirmation blocks: 2
- Block Period: 1
- Gaslimit: 10^9
- DAG BFT (Batch Size: 2MB, 1 worker, 1 primary)
- Send rate: 700




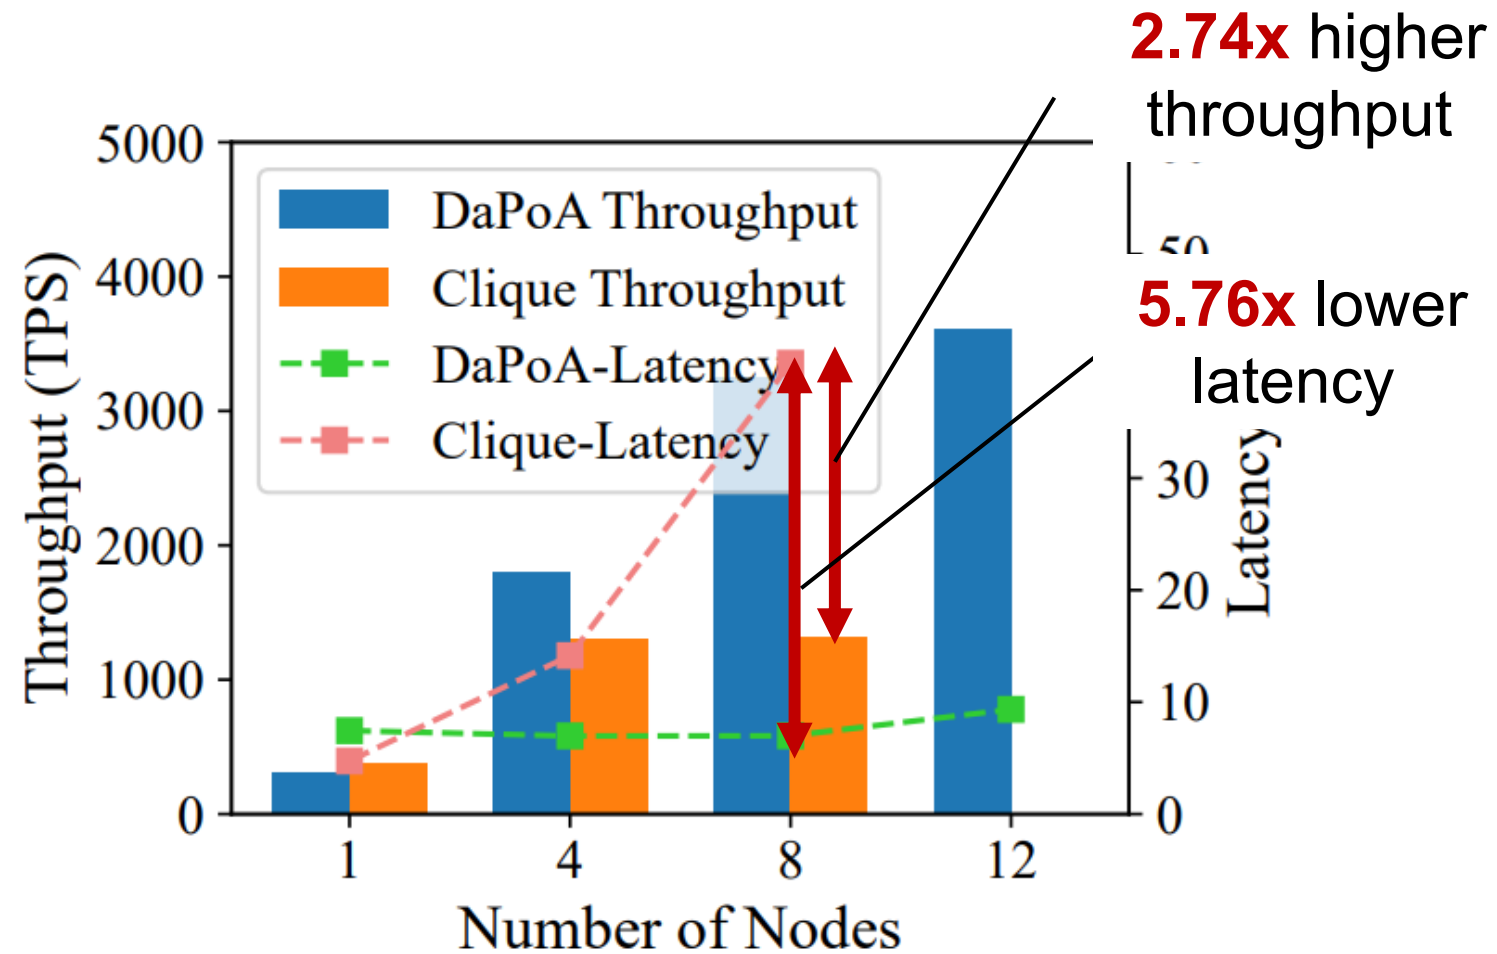
DaPoA scales better

Evaluation: vs. Clique

- Implement DaPoA based on Ethereum Geth*
- Narwhal/Tusk** for DAG BFT
- Benchmark using Hyperledger Caliper
- Hardware: AMD Ryzen 3990X CPU, 256 GB RAM

Parameters

- Simple payment contract
- TX confirmation blocks: 2
- Block Period: 1
- Gaslimit: 10^9
- DAG BFT (Batch Size: 2MB, 1 worker, 1 primary)
- Send rate: 700



DaPoA scales better

Conclusion

DaPoA

- An effort to enhance Ethereum Clique for performance scalability
 - Multiple block acceptance
- Leveraging DAG-based BFT consensus
 - A consistent block view across nodes using DAG mempool.
 - A consistent block ordering using local consensus based on shared randomness